

# Neural Networks & Deep Learning - ICP-6

CS 5720 (CRN 23216)

Student ID: 700745451

Student Name: Kamala Ramesh

1. To implement the use case provided in class – a. To add more Dense layers to the given code and see how accuracy changes.

```
[39] from google.colab import drive
drive.mount('/content/gdrive')

Drive already mounted at /content/gdrive; to attempt to forcibly remount, call drive.mount("/content/gdrive", force_remount=True).

[40] #read the diabetes csv file
path_to_csv = '/content/gdrive/My Drive/diabetes.csv'

import keras
import pandas
from keras.models import Sequential
from keras.layers.core import Dense, Activation

# load dataset
from sklearn.model_selection import train_test_split
import pandas as pd
import numpy as np

dataset = pd.read_csv(path_to_csv, header=None).values

X_train, X_test, Y_train, Y_test = train_test_split(dataset[:,0:8], dataset[:,8],
                                                    test_size=0.25, random_state=87)

np.random.seed(155)
my_first_nn = Sequential() # create model
my_first_nn.add(Dense(20, input_dim=8, activation='relu')) # hidden layer
my_first_nn.add(Dense(1, activation='sigmoid')) # output layer
my_first_nn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
my_first_nn_fitted = my_first_nn.fit(X_train, Y_train, epochs=100,
                                      initial_epoch=0)

print(my_first_nn.summary())
print(my_first_nn.evaluate(X_test, Y_test))

Epoch 1/100
18/18 [=====] - 1s 3ms/step - loss: 6.9816 - acc: 0.6441
Epoch 2/100
18/18 [=====] - 0s 3ms/step - loss: 4.5781 - acc: 0.6458
Epoch 3/100
18/18 [=====] - 0s 4ms/step - loss: 3.6688 - acc: 0.6372
Epoch 4/100
18/18 [=====] - 0s 4ms/step - loss: 3.1913 - acc: 0.6163
```

```
+ Code + Text
Epoch 97/100
18/18 [=====] - 0s 2ms/step - loss: 0.5456 - acc: 0.7413
Epoch 98/100
18/18 [=====] - 0s 2ms/step - loss: 0.5450 - acc: 0.7378
Epoch 99/100
18/18 [=====] - 0s 2ms/step - loss: 0.5612 - acc: 0.7170
Epoch 100/100
18/18 [=====] - 0s 3ms/step - loss: 0.5351 - acc: 0.7483
Model: "sequential_15"

Layer (type)              Output Shape              Param #
=====
dense_34 (Dense)           (None, 20)                180
dense_35 (Dense)           (None, 1)                  21
=====
Total params: 201
Trainable params: 201
Non-trainable params: 0

None
6/6 [=====] - 0s 3ms/step - loss: 0.6220 - acc: 0.6667
[0.6220337748527527, 0.6666666865348816]
```

### After adding dense layers to the existing code

```
{x}
my_second_nn = Sequential() # create model
my_second_nn.add(Dense(20, input_dim=8, activation='relu')) # hidden layer
my_second_nn.add(Dense(10, input_dim=8, activation='relu'))
my_second_nn.add(Dense(5, input_dim=8, activation='relu'))
my_second_nn.add(Dense(1, activation='sigmoid')) # output layer
my_second_nn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
my_second_nn_fitted = my_second_nn.fit(X_train, Y_train, epochs=100,
                                       initial_epoch=0)

print(my_second_nn.summary())
print(my_second_nn.evaluate(X_test, Y_test))

Epoch 82/100
18/18 [=====] - 0s 2ms/step - loss: 0.5229 - acc: 0.7448
Epoch 83/100
18/18 [=====] - 0s 2ms/step - loss: 0.5219 - acc: 0.7431
Epoch 84/100
18/18 [=====] - 0s 4ms/step - loss: 0.5210 - acc: 0.7448
Epoch 85/100
18/18 [=====] - 0s 3ms/step - loss: 0.5229 - acc: 0.7465
Epoch 86/100
18/18 [=====] - 0s 2ms/step - loss: 0.5272 - acc: 0.7396
Epoch 87/100
18/18 [=====] - 0s 3ms/step - loss: 0.5194 - acc: 0.7569
Epoch 88/100
18/18 [=====] - 0s 3ms/step - loss: 0.5172 - acc: 0.7361
Epoch 89/100
18/18 [=====] - 0s 2ms/step - loss: 0.5184 - acc: 0.7448
Epoch 90/100
18/18 [=====] - 0s 2ms/step - loss: 0.5157 - acc: 0.7396
Epoch 91/100
```

```

Epoch 99/100
18/18 [=====] - 0s 2ms/step - loss: 0.5154 - acc: 0.7448
Epoch 100/100
18/18 [=====] - 0s 2ms/step - loss: 0.5097 - acc: 0.7500
Model: "sequential_16"

```

Layer (type)	Output Shape	Param #
dense_36 (Dense)	(None, 20)	180
dense_37 (Dense)	(None, 10)	210
dense_38 (Dense)	(None, 5)	55
dense_39 (Dense)	(None, 1)	6

```

=====
Total params: 451
Trainable params: 451
Non-trainable params: 0
None
6/6 [=====] - 0s 3ms/step - loss: 0.5674 - acc: 0.7031
[0.5673847198486328, 0.703125]

```

For this execution, the loss value decreased, and accuracy increased.

2. Perform the same code as above for the breast cancer dataset and changes the input values as required.

```

+ Code + Text
[1] from google.colab import drive
    drive.mount('/content/gdrive')
Mounted at /content/gdrive
[25] path_to_csv_2 = '/content/gdrive/My Drive/breastcancer.csv'

```

```

import keras
import pandas
from keras.models import Sequential
from keras.layers.core import Dense, Activation

# load dataset
from sklearn.model_selection import train_test_split
import pandas as pd
import numpy as np

dataset = pd.read_csv(path_to_csv_2)

#label the output data from string to numerical value
from sklearn.preprocessing import LabelEncoder

lb_make = LabelEncoder()
dataset["diagnosis_code"] = lb_make.fit_transform(dataset["diagnosis"])
dataset["diagnosis_code"].value_counts()

#Column 'id' is not included in input as it is unique for each record
X = dataset[['radius_mean', 'texture_mean', 'perimeter_mean',
             'area_mean', 'smoothness_mean', 'compactness_mean', 'concavity_mean',
             'concave points_mean', 'symmetry_mean', 'fractal_dimension_mean',
             'radius_se', 'texture_se', 'perimeter_se', 'area_se', 'smoothness_se',
             'compactness_se', 'concavity_se', 'concave points_se', 'symmetry_se',
             'fractal_dimension_se', 'radius_worst', 'texture_worst',
             'perimeter_worst', 'area_worst', 'smoothness_worst',
             'compactness_worst', 'concavity_worst', 'concave points_worst',
             'symmetry_worst', 'fractal_dimension_worst']]

y = dataset['diagnosis_code']

```

```

#Splitting the data
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, shuffle=True, random_state=5)

#training the model
np.random.seed(155)
my_first_nn = Sequential() # create model
my_first_nn.add(Dense(30, input_dim=30, activation='relu')) # hidden layer
my_first_nn.add(Dense(1, activation='sigmoid')) # output layer
my_first_nn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
my_first_nn_fitted = my_first_nn.fit(X_train, y_train, epochs=100,
                                     initial_epoch=0)

print(my_first_nn.summary())
print(my_first_nn.evaluate(X_test, y_test))

```

```

Epoch 1/100
15/15 [=====] - 1s 2ms/step - loss: 78.2931 - acc: 0.6264
Epoch 2/100
15/15 [=====] - 0s 2ms/step - loss: 17.3186 - acc: 0.6615
Epoch 3/100
15/15 [=====] - 0s 2ms/step - loss: 4.2325 - acc: 0.5253
Epoch 4/100
15/15 [=====] - 0s 2ms/step - loss: 0.8222 - acc: 0.8791
Epoch 5/100
15/15 [=====] - 0s 2ms/step - loss: 0.5686 - acc: 0.8593
Epoch 6/100
15/15 [=====] - 0s 2ms/step - loss: 0.4340 - acc: 0.9011
Epoch 7/100
15/15 [=====] - 0s 2ms/step - loss: 0.4172 - acc: 0.8747
Epoch 8/100
15/15 [=====] - 0s 2ms/step - loss: 0.3826 - acc: 0.9143

```

```

Epoch 96/100
15/15 [=====] - 0s 2ms/step - loss: 0.2280 - acc: 0.9209
Epoch 97/100
15/15 [=====] - 0s 2ms/step - loss: 0.1556 - acc: 0.9407
Epoch 98/100
15/15 [=====] - 0s 2ms/step - loss: 0.1606 - acc: 0.9363
Epoch 99/100
15/15 [=====] - 0s 2ms/step - loss: 0.1827 - acc: 0.9297
Epoch 100/100
15/15 [=====] - 0s 2ms/step - loss: 0.1717 - acc: 0.9341
Model: "sequential_10"

Layer (type)                 Output Shape              Param #
=====
dense_20 (Dense)              (None, 30)                930
dense_21 (Dense)              (None, 1)                 31
=====
Total params: 961
Trainable params: 961
Non-trainable params: 0

None
4/4 [=====] - 0s 4ms/step - loss: 0.1465 - acc: 0.9298
[0.1465480476617813, 0.9298245906829834]

```

### 3. Use the StandardScaler to normalize the data and perform the same.

```

[28] #normalize the data using StandardScaler
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
sc.fit(X)
X_scaled = sc.transform(X)

X_train, X_test, y_train, y_test = train_test_split(X_scaled, y, stratify=y, test_size=0.2, shuffle=True, random_state=5)

#train the model using the normalized input set
np.random.seed(155)
my_first_nn = Sequential() # create model
my_first_nn.add(Dense(30, input_dim=30, activation='relu')) # hidden layer
my_first_nn.add(Dense(1, activation='sigmoid')) # output layer
my_first_nn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['acc'])
my_first_nn_fitted = my_first_nn.fit(X_train, y_train, epochs=100,
                                     initial_epoch=0)

print(my_first_nn.summary())
print(my_first_nn.evaluate(X_test, y_test))

Epoch 1/100
15/15 [=====] - 1s 2ms/step - loss: 0.5049 - acc: 0.8198
Epoch 2/100
15/15 [=====] - 0s 2ms/step - loss: 0.3493 - acc: 0.9055
Epoch 3/100
15/15 [=====] - 0s 2ms/step - loss: 0.2645 - acc: 0.9319
Epoch 4/100
15/15 [=====] - 0s 2ms/step - loss: 0.2136 - acc: 0.9385
Epoch 5/100

```

```

Epoch 94/100
15/15 [=====] - 0s 2ms/step - loss: 0.0259 - acc: 0.9912
Epoch 95/100
15/15 [=====] - 0s 3ms/step - loss: 0.0257 - acc: 0.9912
Epoch 96/100
15/15 [=====] - 0s 2ms/step - loss: 0.0253 - acc: 0.9912
Epoch 97/100
15/15 [=====] - 0s 2ms/step - loss: 0.0250 - acc: 0.9912
Epoch 98/100
15/15 [=====] - 0s 2ms/step - loss: 0.0247 - acc: 0.9912
Epoch 99/100
15/15 [=====] - 0s 2ms/step - loss: 0.0243 - acc: 0.9934
Epoch 100/100
15/15 [=====] - 0s 3ms/step - loss: 0.0242 - acc: 0.9934
Model: "sequential_11"

Layer (type)                 Output Shape              Param #
=====
dense_22 (Dense)             (None, 30)                930
dense_23 (Dense)             (None, 1)                 31
=====
Total params: 961
Trainable params: 961
Non-trainable params: 0

None
4/4 [=====] - 0s 4ms/step - loss: 0.0676 - acc: 0.9649
[0.06761771440505981, 0.9649122953414917]

```

When we feed the normalized data to the network the accuracy increases, and the loss decreases.

## Image Classification Problem

```
[57] from keras import Sequential
      from keras.datasets import mnist
      import numpy as np
      from keras.layers import Dense
      from keras.utils import to_categorical

      (train_images, train_labels), (test_images, test_labels) = mnist.load_data()

      print(train_images.shape[1:])
      #process the data
      #1. convert each image of shape 28*28 to 784 dimensional which will be fed to the network as a single feature
      dimData = np.prod(train_images.shape[1:])
      print(dimData)
      train_data = train_images.reshape(train_images.shape[0], dimData)
      test_data = test_images.reshape(test_images.shape[0], dimData)

      #convert data to float and scale values between 0 and 1
      train_data = train_data.astype('float')
      test_data = test_data.astype('float')
      #scale data
      train_data /= 255.0
      test_data /= 255.0
      #change the labels from integer to one-hot encoding. to_categorical is doing the same thing as LabelEncoder()
      train_labels_one_hot = to_categorical(train_labels)
      test_labels_one_hot = to_categorical(test_labels)


      #creating network
      model = Sequential()
      model.add(Dense(512, activation='relu', input_shape=(dimData,)))
      model.add(Dense(512, activation='relu'))
      model.add(Dense(10, activation='softmax'))

      model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
      history = model.fit(train_data, train_labels_one_hot, batch_size=256, epochs=10, verbose=1,
                          validation_data=(test_data, test_labels_one_hot))


      (28, 28)
      784
      Epoch 1/10
      235/235 [=====] - 9s 36ms/step - loss: 0.2906 - accuracy: 0.9115 - val_loss: 0.1891 - val_accuracy: 0.9377
      Epoch 2/10
      235/235 [=====] - 7s 30ms/step - loss: 0.1021 - accuracy: 0.9683 - val_loss: 0.1172 - val_accuracy: 0.9609
      Epoch 3/10
      235/235 [=====] - 8s 34ms/step - loss: 0.0631 - accuracy: 0.9804 - val_loss: 0.0773 - val_accuracy: 0.9752
      Epoch 4/10
      235/235 [=====] - 8s 33ms/step - loss: 0.0448 - accuracy: 0.9858 - val_loss: 0.0811 - val_accuracy: 0.9751
      Epoch 5/10
      235/235 [=====] - 7s 32ms/step - loss: 0.0318 - accuracy: 0.9894 - val_loss: 0.0750 - val_accuracy: 0.9789
      Epoch 6/10
      235/235 [=====] - 8s 35ms/step - loss: 0.0247 - accuracy: 0.9922 - val_loss: 0.0697 - val_accuracy: 0.9809
      Epoch 7/10
      235/235 [=====] - 7s 29ms/step - loss: 0.0168 - accuracy: 0.9947 - val_loss: 0.0710 - val_accuracy: 0.9798
      Epoch 8/10
      235/235 [=====] - 8s 33ms/step - loss: 0.0135 - accuracy: 0.9959 - val_loss: 0.0661 - val_accuracy: 0.9821
      Epoch 9/10
      235/235 [=====] - 7s 29ms/step - loss: 0.0094 - accuracy: 0.9975 - val_loss: 0.0711 - val_accuracy: 0.9823
      Epoch 10/10
      235/235 [=====] - 8s 33ms/step - loss: 0.0075 - accuracy: 0.9976 - val_loss: 0.0944 - val_accuracy: 0.9792
```

1. Plot the loss and accuracy for both training data and validation data using the history object in the source code

```
✓ 0s ▶ import matplotlib.pyplot as plt
# Plot the training and validation accuracy over epochs
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



```
✓ 0s [63] # Plot the training and validation loss over epochs
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```





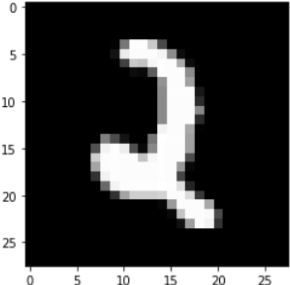
2. Plot one of the images in the test data, and then do inferencing to check what is the prediction of the model on that single image.

```
0s # Select a random image from the test data
idx = np.random.randint(0, test_images.shape[0])
img = test_images[idx]

# Plot the selected image
plt.imshow(img, cmap='gray')
plt.show()

input_image = img.reshape(1, 784).astype('float32') / 255.0

prediction = model.predict(input_image)
print('The image is predicted as:', np.argmax(prediction))
```



```
1/1 [=====] - 0s 26ms/step
The image is predicted as: 2
```

3. To change the number of hidden layer and the activation to tanh or sigmoid and observe the changes

```
1m #creating network
model_upd = Sequential()
model_upd.add(Dense(512, activation='tanh', input_shape=(dimData,)))
model_upd.add(Dense(512, activation='tanh'))
model_upd.add(Dense(256, activation='tanh'))
model_upd.add(Dense(128, activation='tanh'))
model_upd.add(Dense(10, activation='softmax'))

model_upd.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_upd.fit(train_data, train_labels_one_hot, batch_size=256, epochs=10, verbose=1,
                        validation_data=(test_data, test_labels_one_hot))
```

```
Epoch 1/10
235/235 [=====] - 10s 41ms/step - loss: 0.3548 - accuracy: 0.8903 - val_loss: 0.1926 - val_accuracy: 0.9444
Epoch 2/10
235/235 [=====] - 8s 36ms/step - loss: 0.1455 - accuracy: 0.9558 - val_loss: 0.1749 - val_accuracy: 0.9456
Epoch 3/10
235/235 [=====] - 10s 42ms/step - loss: 0.0980 - accuracy: 0.9694 - val_loss: 0.1277 - val_accuracy: 0.9602
Epoch 4/10
235/235 [=====] - 10s 42ms/step - loss: 0.0705 - accuracy: 0.9779 - val_loss: 0.2338 - val_accuracy: 0.9230
Epoch 5/10
235/235 [=====] - 9s 36ms/step - loss: 0.0543 - accuracy: 0.9831 - val_loss: 0.0902 - val_accuracy: 0.9726
Epoch 6/10
235/235 [=====] - 10s 42ms/step - loss: 0.0396 - accuracy: 0.9876 - val_loss: 0.1559 - val_accuracy: 0.9531
Epoch 7/10
235/235 [=====] - 10s 42ms/step - loss: 0.0304 - accuracy: 0.9903 - val_loss: 0.1162 - val_accuracy: 0.9650
Epoch 8/10
235/235 [=====] - 9s 40ms/step - loss: 0.0236 - accuracy: 0.9929 - val_loss: 0.0875 - val_accuracy: 0.9727
Epoch 9/10
235/235 [=====] - 8s 36ms/step - loss: 0.0184 - accuracy: 0.9943 - val_loss: 0.0765 - val_accuracy: 0.9765
Epoch 10/10
235/235 [=====] - 10s 42ms/step - loss: 0.0135 - accuracy: 0.9959 - val_loss: 0.0759 - val_accuracy: 0.9785
```

There is a very minute changes when we add more dense layers and change the activation to function to tanh. Here for the validation data, the loss decreased and the accuracy also decreased.

#### 4. Executing the same code withing scaling the input data

```
from keras import Sequential
from keras.datasets import mnist
import numpy as np
from keras.layers import Dense
from keras.utils import to_categorical

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()

print(train_images.shape[1:])
#process the data
#1. convert each image of shape 28*28 to 784 dimensional which will be fed to the network as a single feature
dimData = np.prod(train_images.shape[1:])
print(dimData)
train_data = train_images.reshape(train_images.shape[0], dimData)
test_data = test_images.reshape(test_images.shape[0], dimData)

#convert data to float and scale values between 0 and 1
train_data = train_data.astype('float')
test_data = test_data.astype('float')

#Commenting the scale data part
#train_data /=255.0
#test_data /=255.0

#change the labels from integer to one-hot encoding. to_categorical is doing the same thing as LabelEncoder()
train_labels_one_hot = to_categorical(train_labels)
test_labels_one_hot = to_categorical(test_labels)

#creating network
model = Sequential()
model.add(Dense(512, activation='relu', input_shape=(dimData,)))
model.add(Dense(512, activation='relu'))
model.add(Dense(10, activation='softmax'))

#Feeding the unscaled data to the network
model.compile(optimizer='rmsprop', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(train_data, train_labels_one_hot, batch_size=256, epochs=10, verbose=1,
                    validation_data=(test_data, test_labels_one_hot))
```

Downloading data from <https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz>  
11490434/11490434 [=====] - 1s 0us/step  
(28, 28)  
784  
Epoch 1/10  
235/235 [=====] - 13s 50ms/step - loss: 7.1164 - accuracy: 0.8726 - val\_loss: 0.7212 - val\_accuracy: 0.9056  
Epoch 2/10  
235/235 [=====] - 5s 20ms/step - loss: 0.4010 - accuracy: 0.9453 - val\_loss: 0.6054 - val\_accuracy: 0.9258  
Epoch 3/10  
235/235 [=====] - 5s 21ms/step - loss: 0.2324 - accuracy: 0.9601 - val\_loss: 0.4047 - val\_accuracy: 0.9423  
Epoch 4/10  
235/235 [=====] - 5s 21ms/step - loss: 0.1761 - accuracy: 0.9679 - val\_loss: 0.2555 - val\_accuracy: 0.9577  
Epoch 5/10  
235/235 [=====] - 5s 20ms/step - loss: 0.1583 - accuracy: 0.9727 - val\_loss: 0.2695 - val\_accuracy: 0.9648  
Epoch 6/10  
235/235 [=====] - 5s 22ms/step - loss: 0.1351 - accuracy: 0.9761 - val\_loss: 0.2610 - val\_accuracy: 0.9659  
Epoch 7/10  
235/235 [=====] - 4s 19ms/step - loss: 0.1237 - accuracy: 0.9801 - val\_loss: 0.2888 - val\_accuracy: 0.9681  
Epoch 8/10  
235/235 [=====] - 5s 21ms/step - loss: 0.1173 - accuracy: 0.9807 - val\_loss: 0.2874 - val\_accuracy: 0.9691  
Epoch 9/10  
235/235 [=====] - 6s 26ms/step - loss: 0.1077 - accuracy: 0.9829 - val\_loss: 0.2675 - val\_accuracy: 0.9734  
Epoch 10/10  
235/235 [=====] - 6s 24ms/step - loss: 0.1053 - accuracy: 0.9854 - val\_loss: 0.3169 - val\_accuracy: 0.9707

When we feed the network without scaling the data, the loss increases and the accuracy decreases.